

# HPC for Dense Stellar Systems

Cristián Maureira-Fredes

Albert Einstein Institute

October 19, 2017

# Contents

## Introduction

The  $N$ -body problem

Computational aspects

CPU approach

GPU Computing

Programming strategy

## Experiments

## Projects

## Future work

# The $N$ -body problem

## Definition (1/2)

Predicting motion of a group or celestial objects that interact each other gravitationally.




Figure: M69 in Sagittarius.

# The $N$ -body problem

## Definition (2/2)

Purely dynamic problem, in which the bodies orbital evolution is determined exclusive by the **gravitational interaction**,

$$\ddot{\vec{r}}_i = -G \sum_{\substack{j=1 \\ j \neq i}}^N m_j \frac{\vec{r}_{ij}}{|\vec{r}_{ij}|^3}, \quad (1)$$

where  $G$  is the gravitational constant,  $m_j$  is the mass of the  $j$ -th particle and  $\vec{r}_{ij} = (\vec{r}_i - \vec{r}_j)$  the position in *Cartesian* coordinates.

# The $N$ -body problem

## Checking the system evolution

- ▶ The **initial condition** are usually the masses, position and velocity. (Different distributions and shapes: Plummer, King, Dehnen, etc)
- ▶ **Chaotic nature**, the evolution of the system is highly sensitive to the initial conditions.
- ▶ The often invariant to check the integration of the system, is the system's **energy**,

$$E = K + U \quad (2)$$

and sometimes the **angular momentum**,

$$\vec{L} = \vec{r} \times \vec{p} \quad (3)$$

# The $N$ -body problem

## $N$ -body algorithms classification

**Collision-less** A star just sees the **background potential** of the rest of the stellar system. A model of this situation is the Barnes-Hut Treecode with a complexity  $O(N \log N)$  [1] or the fast multipole method with  $O(N)$  [2].

**Collisional ("direct-summation")** One star integrates **all gravitational forces** for all stars. This typically scale as  $O(N^2)$ . A well-known example is the family of algorithm of Aarseth the direct-summation NBODY integrator [3, 4, 5] or KIRA code [6].

# The $N$ -body problem

## Moving the particles (Timesteps) (1/3)

- ▶ Individual timesteps allows an accurate treatment of the evolution (handle close encounters)
- ▶ An Predictor-Corrector integration scheme:
  1. **Select** a particle  $i$  which has the minimum  $\Delta t_i + t_i$  (where  $\Delta t_i$  and  $t_i$  are its own timestep and current time)
  2. Calculate gravitational **interactions** of this particle. (Predict all the other particles to the current time)
  3. **Integrate** the particle  $i$  to its new time  $t_{new} = \Delta t_i + t_i$ .  
(Correcting its information using higher derivatives elements)
  4. Get the new timestep.
- ▶ **Difficult** scenario for parallel computing.

# The $N$ -body problem


## Moving the particles (Timesteps) (2/3)

- ▶ Block timesteps offers a **restriction** for the particles timesteps, to be a power of two [7].
- ▶ The integration scheme will change from "Select a particle", for "Select a **group** of particles".
- ▶ This time step scheme is popular among  $N$ -body code, like Starlab [8, 9], Aarseth  $N$ -body codes [3, 5, 10],  $\phi$ GRAPE [11].



# The $N$ -body problem

### Moving the particles (Timesteps) (3/3)



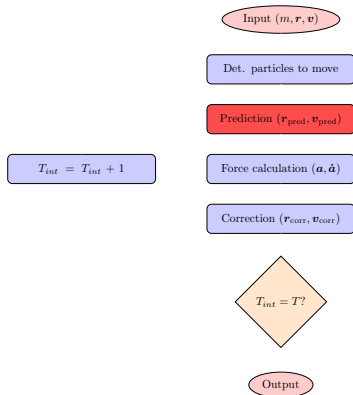
**Figure:** Block time steps illustration. The different blocks are represented by different colors. Each particle is predicted (not move) at every time  $t$  (gray arrows), even if it's not their block time-step (blue circles). The particles will be updated (moved) only in their block time-step (black circles).

# The $N$ -body problem

## Hermite 4th order (Predictor-Corrector)

### Prediction

$$\vec{r}_{i,pred} = \vec{r}_{i,0} + \vec{v}_{i,0}\Delta t_i + \vec{a}_{i,0}\frac{\Delta t_i^2}{2!} + \vec{\ddot{a}}_{i,0}\frac{\Delta t_i^3}{3!}$$
$$\vec{v}_{i,pred} = \vec{v}_{i,0} + \vec{a}_{i,0}\Delta t_i + \vec{\dot{a}}_{i,0}\frac{\Delta t_i^2}{2!}$$



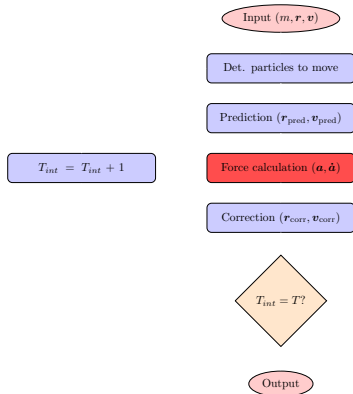
# The $N$ -body problem

Hermite 4th order (Predictor-Corrector)

## Force calculation

$$\vec{a}_{i,1} = \sum_{\substack{j=0 \\ j \neq i}}^N Gm_j \frac{\vec{r}_{ij}}{(r_{ij}^2 + \epsilon^2)^{\frac{3}{2}}},$$

$$\vec{a}_{i,1} = \sum_{\substack{j=0 \\ j \neq i}}^N Gm_j \left[ \frac{\vec{v}_{ij}}{(r_{ij}^2 + \epsilon^2)^{\frac{3}{2}}} - \frac{3(\vec{v}_{ij} \cdot \vec{r}_{ij})\vec{r}_i}{(r_{ij}^2 + \epsilon^2)^{\frac{5}{2}}} \right],$$

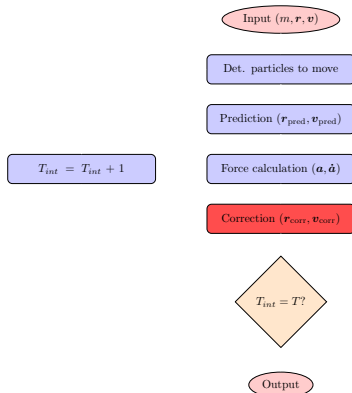


# The $N$ -body problem

## Hermite 4th order (Predictor-Corrector)

### Correction

$$\begin{aligned}\vec{r}_{i,1} &= \vec{r}_{i,pred} + \frac{1}{24}\Delta t_i^4 \vec{a}_{i,0}^{(2)} + \frac{1}{120}\Delta t_i^5 \vec{a}_{i,0}^{(3)} \\ \vec{v}_{i,1} &= \vec{v}_{i,pred} + \frac{1}{4}\Delta t_i^3 \vec{a}_{i,0}^{(2)} + \frac{1}{24}\Delta t_i^4 \vec{a}_{i,0}^{(3)}\end{aligned}$$



# Computational Aspects

## The computational challenge

- ▶ The  $N$ -body codes evolution is related to the available **hardware** in our time.
- ▶ The algorithms with a complexity of  $O(N^2)$  or  $O(N^3)$  require **supercomputers**.
  - ▶ e.g. **beowulf clusters**, which require a parallelization of the code (NBODY6++ developed by Spurzem et al. [4]).
  - ▶ Special-purpose hardware, like the **GRAPE** (short for GRAvity PipE system [12, 13, 14, 15].
- ▶ The literature overview reveals a strong interest on porting the existing codes to the **GPU** architecture, like e.g. the work of [16, 17, 18] on single nodes or using large clusters [19, 10, 20].

# Computational Aspects

## Parallel CPU implementation

- ▶ Single-core implementation to perform a profiling (gprof).
  - ▶ Gravitational interaction is the bottleneck.
  - ▶ Usually  $N_{act} \ll N$ .
- ▶ Many-core with OpenMP.
  - ▶ `#pragma omp parallel for`
- ▶ Many-core with MPI (two implementations)
  - ▶ `MPI_Allreduce`, `MPI_Bcast`

# Computational Aspects

## Parallel CPU implementation

**Listing 1:** Reduce every  $N_{act}$  particle

```
1  for (int i = 0; i < Nact; i++)
2  {
3      ...
4      for (int j = 0; j < N; j++)
5      {
6          gravitational_interaction(...)
              ;
7      }
8      ...
9      MPI_Allreduce(...);
10 }
```

**Listing 2:** Reduce all  $N_{act}$  particle

```
1  for (int i = 0; i < Nact; i++)
2  {
3      ...
4      for (int j = 0; j < N; j++)
5      {
6          gravitational_interaction(...)
              ;
7      }
8      ...
9  }
10 MPI_Allreduce(...);
```

# Computational Aspects

## GPU Computing

*"Using a GPU (Graphic Processing Unit) together with a CPU to accelerate scientific calculation operations or general purpose calculation"*



Figure: NVIDIA® GTX Titan



- ▶ CPU,
  - ▶ Designed to have a good **performance** in parallel and non-parallel scenarios.
  - ▶ Minimizes the **latency** experienced by a thread (large cache memory)
- ▶ GPU,
  - ▶ Designed to perform highly parallel work.
  - ▶ Maximizes the **throughput** of all the threads.

## Performance

Capacity of perform individual instructions in a certain time.

## Latency

Measure of time delay experienced in a system.

## Throughput

Capacity of perform a whole task in a certain time.

# GPU Architecture

## Computational Aspects

### Task parallelism

Each processor perform a different task.

### Data parallelism

Each processor perform the same task, but not on the same data set.




Figure: GPU and CPU core scheme

# Computational Aspects

## Programming strategy





Figure: CUDA Programming strategy

1. CPU memory allocation,
2. GPU memory allocation,
3. Data copying, CPU → GPU,
4. Task execution on the data,
5. Data copying, GPU → CPU,

# Parallelization scheme

## $j$ -parallelization scheme

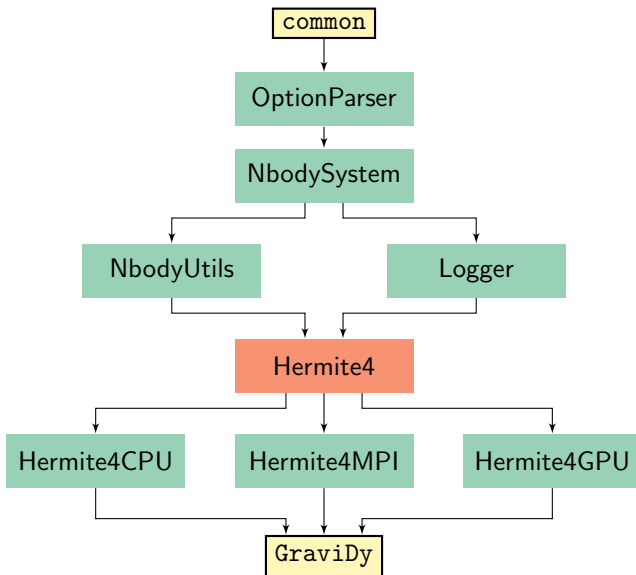
Our configuration is based in the idea presented in [10],



**Figure:** Parallelization scheme to split the  $j$ -loop instead of the  $i$ -loop. In this case, we have two sections, the first is to calculate the force interactions of the  $i$ -particle with the whole system but by different threads. Then a reduction (sum) is necessary to get the new value for the  $i$ -particle force.

# Implementation

## Class diagram



# Experiments

## Details of the hardware


---

CPU	Intel(R) Xeon(R) CPU E5-4650 0 @ 2.70GHz
GPU	Tesla M2050 @ 575 Mhz (448 cores).
RAM	24 GB
OS	Scientific Linux release 6.4

---

# Experiments

## Integrator scaling



**Figure:** Clock time of integration from  $t = 1$  to  $t = 2$  NBU using  $\eta = 0.01$  and  $\epsilon = 10^{-4}$  using different amount of particles.

# Experiments

## Clock time comparison

<b>N</b>	<b>CPU</b>	<b>OpenMP</b>	<b>CPU + GPU</b>	<b>MPI-1</b>	<b>MPI-2</b>	<b>GPU</b>
<b>1k</b>	12	8	3	6	2	1
<b>2k</b>	61	34	13	14	7	3
<b>4k</b>	282	162	54	51	27	9
<b>8k</b>	1227	682	208	105	64	23
<b>16k</b>	5542	3227	904	364	317	82
<b>32k</b>	26383	15076	3722	1247	1145	275

**Table:** Clock time foreach integrator version (in sec).



# Experiments

## Clock time comparison





Figure: Acceleration between the implementations described in Table 1

# Experiments

## Integrator Performance



**Figure:** GPU gravitational interactions performance in GFLOPS for different amount of particles.

# Projects

## Software

The current version of our new  $N$ -body code, written in C/C++ and CUDA, called GRAVIDY .

- ▶ The current version of our code,
  - ▶ Using [Hermite 4th order](#) integration scheme.
  - ▶ Block timesteps for helping parallelism.
  - ▶ [Suitable](#) in the energy conservation, reaching errors around  $\approx 10^{-9}$  and  $\approx 10^{-7}$ .
  - ▶ OO.
  - ▶ Documentation ([Doxygen](#)).

# Projects

## Software development

- ▶ Main software
  - ▶ Main goal in the development, **legibility**.
    - ▶ Easy to read, modify and understand,
  - ▶ **Balance** between optimization and maintainability.
- ▶ Utility scripts.

# Projects

## Milestones

- ▶ Programming

- ▶ Physics

# Projects

## Milestones

- ▶ Programming
  - ▶ Multi-GPU environment, large clusters (MPI).
- ▶ Physics

# Projects

## Milestones

- ▶ Programming
  - ▶ Multi-GPU environment, large clusters (MPI).
  - ▶ Assess, Parallelize, Optimize, Deploy(APOD) design cycle.
- ▶ Physics

# Projects

## Milestones

- ▶ Programming
  - ▶ Multi-GPU environment, large clusters (MPI).
  - ▶ Assess, Parallelize, Optimize, Deploy(APOD) design cycle.
  - ▶ Analysis of FP precision.
- ▶ Physics



# Projects

## Milestones

- ▶ Programming
  - ▶ Multi-GPU environment, large clusters (MPI).
  - ▶ Assess, Parallelize, Optimize, Deploy (APOD) design cycle.
  - ▶ Analysis of FP precision.
- ▶ Physics
  - ▶ Semi-Keplerian systems
    - ▶ Globular cluster + BH.

# Projects

## Milestones

- ▶ Programming
  - ▶ Multi-GPU environment, large clusters (MPI).
  - ▶ Assess, Parallelize, Optimize, Deploy (APOD) design cycle.
  - ▶ Analysis of FP precision.
- ▶ Physics
  - ▶ Semi-Keplerian systems
    - ▶ Globular cluster + BH.
  - ▶ Neighbors treatment.
    - ▶ Near particles are more important than the rest of the system.

# Projects

## Milestones

- ▶ Programming
  - ▶ Multi-GPU environment, large clusters (MPI).
  - ▶ Assess, Parallelize, Optimize, Deploy (APOD) design cycle.
  - ▶ Analysis of FP precision.
- ▶ Physics
  - ▶ Semi-Keplerian systems
    - ▶ Globular cluster + BH.
  - ▶ Neighbors treatment.
    - ▶ Near particles are more important than the rest of the system.
  - ▶ Regularisations (KS, Chain, ...)
    - ▶ Remove the softening parameter  $\epsilon^2$ .

# Projects

## Milestones

- ▶ Programming
  - ▶ Multi-GPU environment, large clusters (MPI).
  - ▶ Assess, Parallelize, Optimize, Deploy(APOD) design cycle.
  - ▶ Analysis of FP precision.
- ▶ Physics
  - ▶ Semi-Keplerian systems
    - ▶ Globular cluster + BH.
  - ▶ Neighbors treatment.
    - ▶ Near particles are more important than the rest of the system.
  - ▶ Regularisations (KS, Chain, ...)
    - ▶ Remove the softening parameter  $\epsilon^2$ .
  - ▶ Higher order integration schemes.
    - ▶ Hermite 6th order (Faster, and more accurate).

# Projects

## Milestones

- ▶ Programming
  - ▶ Multi-GPU environment, large clusters (MPI).
  - ▶ Assess, Parallelize, Optimize, Deploy (APOD) design cycle.
  - ▶ Analysis of FP precision.
- ▶ Physics
  - ▶ Semi-Keplerian systems
    - ▶ Globular cluster + BH.
  - ▶ Neighbors treatment.
    - ▶ Near particles are more important than the rest of the system.
  - ▶ Regularisations (KS, Chain, ...)
    - ▶ Remove the softening parameter  $\epsilon^2$ .
  - ▶ Higher order integration schemes.
    - ▶ Hermite 6th order (Faster, and more accurate).
  - ▶ Smoothed-particle hydrodynamics (SPH).
    - ▶ Different treatment for particle systems.

# Projects

## Milestones

- ▶ Programming
  - ▶ Multi-GPU environment, large clusters (MPI).
  - ▶ Assess, Parallelize, Optimize, Deploy(APOD) design cycle.
  - ▶ Analysis of FP precision.
- ▶ Physics
  - ▶ Semi-Keplerian systems
    - ▶ Globular cluster + BH.
  - ▶ Neighbors treatment.
    - ▶ Near particles are more important than the rest of the system.
  - ▶ Regularisations (KS, Chain, ...)
    - ▶ Remove the softening parameter  $\epsilon^2$ .
  - ▶ Higher order integration schemes.
    - ▶ Hermite 6th order (Faster, and more accurate).
  - ▶ Smoothed-particle hydrodynamics (SPH).
    - ▶ Different treatment for particle systems.
  - ▶ etc.

# HPC for Dense Stellar Systems

Cristián Maureira-Fredes

Albert Einstein Institute

October 19, 2017

-  J. Barnes and P. Hut, "A hierarchical  $O(N \log N)$  force-calculation algorithm," vol. 324, pp. 446–449, Dec. 1986.
-  L. Greendard, *The rapid evaluation of potential fields in particles systems*. PhD thesis, Yale University, New Haven, CT, 1987.
-  S. J. Aarseth, "From NBODY1 to NBODY6: The Growth of an Industry," *The Publications of the Astronomical Society of the Pacific*, vol. 111, pp. 1333–1346, Nov. 1999.
-  R. Spurzem, "Direct N-body Simulations," *Journal of Computational and Applied Mathematics*, vol. 109, pp. 407–432, Sept. 1999.
-  S. J. Aarseth, *Gravitational N-Body Simulations*. ISBN 0521432723. Cambridge, UK: Cambridge University Press, November 2003., Nov. 2003.
-  S. F. Portegies Zwart, S. L. W. McMillan, P. Hut, and J. Makino, "Star cluster ecology - IV. Dissection of an open star cluster: photometry," vol. 321, pp. 199–226, Feb. 2001.
-  W. H. Press, "Techniques and tricks for  $n$ -body computation," in *The Use of Supercomputers in Stellar Dynamics* (P. Hut and S. L. W. McMillan, eds.), p. 184, Springer-Verlag, 1986.





S. F. Portegies Zwart, S. L. W. McMillan, P. Hut, and J. Makino, "Star cluster ecology - IV. Dissection of an open star cluster: photometry," vol. 321, pp. 199–226, Feb. 2001.



P. Hut, "The Starlab Environment for Dense Stellar Systems," in *Astrophysical Supercomputing using Particle Simulations* (J. Makino and P. Hut, eds.), vol. 208 of *IAU Symposium*, p. 331, 2003.



K. Nitadori and S. J. Aarseth, "Accelerating NBODY6 with graphics processing units," vol. 424, pp. 545–552, July 2012.



S. Harfst, A. Gualandris, D. Merritt, and S. Mikkola, "A hybrid N-body code incorporating algorithmic regularization and post-Newtonian forces," vol. 389, pp. 2–12, Sept. 2008.



M. Taiji, J. Makino, T. Fukushige, T. Ebisuzaki, and D. Sugimoto, "Grape-4: A teraflops machine for n-body simulations," in *IAU Symp. 174: Dynamical Evolution of Star Clusters: Confrontation of Theory and Observations* (P. Hut and J. Makino, eds.), p. 141, 1996.



J. Makino and M. Taiji, *Scientific simulations with special-purpose computers : The GRAPE systems*.

Scientific simulations with special-purpose computers : The GRAPE systems /by Junichiro Makino & Makoto Taiji. Chichester ; Toronto : John Wiley & Sons, c1998., 1998.



J. Makino, "Grape-6," *Highlights in Astronomy*, vol. 11, p. 597, 1998.



T. Fukushige, J. Makino, and A. Kawai, "GRAPE-6A: A Single-Card GRAPE-6 for Parallel PC-GRAPE Cluster Systems," vol. 57, pp. 1009–1021, Dec. 2005.



S. F. Portegies Zwart, R. G. Belleman, and P. M. Geldof, "High-performance direct gravitational N-body simulations on graphics processing units," vol. 12, pp. 641–650, Nov. 2007.



T. Hamada and T. Iitaka, "The Chamomile Scheme: An Optimized Algorithm for N-body simulations on Programmable Graphics Processing Units," Mar. 2007.



R. G. Belleman, J. Bédorf, and S. F. Portegies Zwart, "High performance direct gravitational N-body simulations on graphics processing units II: An implementation in CUDA," vol. 13, pp. 103–112, Feb. 2008.



P. Berczik, K. Nitadori, S. Zhong, R. Spurzem, T. Hamada, X. Wang, I. Berentzen, A. Veles, and W. Ge, "High performance massively parallel direct n-body simulations on large gpu clusters," 2011.



R. Capuzzo-Dolcetta, M. Spera, and D. Punzo, "A fully parallel, high precision, N-body code running on hybrid computing platforms," *Journal of Computational Physics*, vol. 236, pp. 580–593, Mar. 2013.